

ASUTOSH PANDA

📍 Brahmapur, India • ✉ asutosh.pda@gmail.com •  <https://linkedin.com/in/75asu> •  <https://75asu.pages.dev> •  <https://github.com/75asu>

SKILLS

GPU & AI Infrastructure: NVIDIA GPU Operator, DCGM, drivers/CUDA, vLLM, llm-d, KubeRay/Ray Serve, LiteLLM, APISIX, Envoy AI Gateway, Gateway API Inference Extension, NIXL, P/D disaggregation, prefix-aware routing, Karpenter GPU autoscaling, inference serving (TTFT/TPOT/goodput, cost per 1K tokens), guidellm benchmarking, GPU scheduling, AWS Trainium, Amazon Bedrock, SageMaker, Vertex AI

Kubernetes & Infrastructure: Kubernetes (GKE, EKS, bare-metal / data centre), Helm, Kustomize, FluxCD, ArgoCD, Terraform/OpenTofu, Terragrunt, Ansible, Karpenter, AWS Transit Gateway, Istio, Linkerd, external-dns, Bottlerocket, GCP, AWS, Cloudflare, Hetzner, OVH

Observability & Reliability: Datadog, Prometheus, Grafana, Mimir, Grafana Alloy, OpenTelemetry/OTLP, Loki, Tempo, Jaeger, Perses, VictoriaMetrics, VictoriaLogs, VictoriaTraces, vmalert, Alertmanager, SLOs and error budgets, on-call and incident response

CI/CD, Secrets and Access: GitHub Actions, self-hosted ARC runners, trunk-based release and rollback; Workload Identity Federation, GHE OIDC, IRSA + EKS Pod Identity, External Secrets Operator, cert-manager, Pomerium, Cloudflare Zero Trust, GKE Connect Gateway

Development & Data: Golang, Python, Bash, PostgreSQL, CloudNativePG, Aurora, ElastiCache, Starrocks, Redis, DMS, Pub/Sub, Linux

PROFESSIONAL EXPERIENCE

Socure

Chennai, India

Senior Site Reliability Engineer (Remote)

August 2026 - Present

- Own the Fravity to Socure platform migration end to end: architected the target state, sequenced it across **5** teams, and cleared the account-boundary, PII and security sign-offs each phase depended on, with no release freeze for the teams shipping through it.
- Built the GPU inference platform on EKS: Karpenter-scaled scale-from-zero L4oS running vLLM on Ray Serve, Envoy AI Gateway plus Gateway API Inference Extension dispatching to llm-d, APISIX auth and rate limits, benchmarked to **~527K** requests/day at FP8.
- Made inference performance measurable: guidellm precision and traffic sweeps to **p999** against chat-tier SLOs, an OOM crashloop root-caused at **1,400+** restarts, and GPUs halved by correcting an fp32 assumption to bf16.
- Landed the whole platform on AWS across **104** merged PRs in **11** repos, replacing every GCP dependency with a managed equivalent: Cloud SQL to Aurora PostgreSQL, Pub/Sub to SQS, Workload Identity to IRSA; migrated **70** repos to GHE in **< 5** days.
- Took **4** services live end to end, then closed the defect class green deploys could not detect: config falling back to shipped defaults.
- Shipped per-PR ephemeral environments for **4** components that need no cluster access and no SRE in the loop, scoped by an IRSA role granting inference only so unreviewed branches reach no storage, queue or secret; a **15**-hour lockout class now self-heals in **5** minutes.
- Own the observability contract for a platform with no agent and no ticket to add one: aligned it across platform SRE and two product teams, then shipped Datadog to **4** environments, ArgoCD in non-prod and Terraform in prod, **51** monitors, GPU tier on DCGM and vLLM latency.
- Deployed the **3**-environment EKS platform on /21 IPAM VPCs with TGW hub-and-spoke, then set the golden path: keyless GHE OIDC to ECR behind **4** scanners, Istio mTLS, FIPS golden bases, and **10** reusable Actions for trunk-based promotion with hotfix lanes.

ONE2N

Pune, Maharashtra, India

SRE Lead - Client: Fravity AI (Remote)

May 2026 - July 2026

- Built the org's first observability stack, then re-platformed it off AGPL onto VictoriaMetrics on a founder licence call: a dual-run Alloy fleet moved **~305K** series from **192** targets in **4** clusters with no data-loss window, and metrics, logs, traces and alerts behind one Perses view.
- Built fravity-spine, the internal Kubernetes platform other teams consume: private regional GKE, CMEK etcd, Workload Identity, CloudNativePG, OpenTofu across **6** state layers, and identity-based kubectrl retiring bastions and shared SSH.
- Led cost + access governance across **5** GCP + **5** AWS accounts as lead of a two-person SRE team: surfaced **19** live SA keys and **~\$26,520/yr** of list-price exposure, then shipped an org-wide FinOps board and IAM as Terraform on keyless CI.

LUXOR TECHNOLOGY

Seattle, WA, US

DevOps Engineer (Remote)

June 2023 - April 2026

- Operated a multi-cluster Kubernetes platform for **20+** teams and **3,000+** enterprise customers: **8+** clusters across Hetzner, OVH, and GCP with **24/7** uptime for L1/L2 validator and RPC workloads, including **2** Hetzner bare-metal HA clusters (KubeOne, OpenEBS) built from scratch. Standardised on a 7-layer FluxCD GitOps pipeline with Linkerd mTLS, migrating **70+** secrets to Workload Identity.
- Owned CI/CD for **50+** repositories on self-hosted ARC runners. Ported a failing Playwright E2E suite from Mac hosted runners (**\$5,000/mo**) to Linux ARC, cutting runtime **40 to 15 min** and cost to **\$600/mo**. Built code-signed Windows and Mac release pipelines on FluxCD.

TATA CONSULTANCY SERVICES

Hyderabad, Telangana, India

Systems Engineer (Remote)

October 2020 - June 2023

- Migrated **2,500+** repositories (GitLab, Bitbucket, SVN) to GitHub Enterprise for **120** teams and built their CI/CD pipelines, resolving self-hosted-API pagination limits across the platforms and saving **15 weeks** of manual effort.
- Built PowerShell + Python automation (server monitoring, alerting, log rotation, Oracle DB replication, MS SQL updates) across **22** RHEL/WebLogic servers and **3** projects, saving **380+ man-hours/month**.
- Shipped a Golang CLI that cut HP storage provisioning from **30 man-hours** across **30** monthly tickets to a single command, adopted by the storage team as their primary provisioning interface.

OPEN-SOURCE PROJECTS

- ducat** (Jun 2026 - Present)  - **Prometheus** cost exporter in Python: cloud, SaaS + AI spend in one Grafana view (**FOCUS** spec).
- Frit** (May 2026 - Present)  - GPU SRE lab: **vLLM** + **Ray Serve** on GPU Operator; inference-perf (**TTFT/TPOT**, DCGM) in **Grafana**.
- Platform Zero** (Apr 2026 - Present)  - **4** SRE labs (k8slab, ansiblelab, awslab, gcplab); **0-to-100** in one command.
- Truss** (Jun 2026 - Present)  - Self-hostable **Supabase/Appwrite** alternative, with an AI-agent **MCP server**.
- florin** (Jun 2026 - Present)  - Portable DB workbench (VS Code): **Postgres** + **Redis**, connections + queries synced to a **git vault** you own.

TEACHING EXPERIENCE

- Full-Stack Instructor, **4** cohorts (2022-2026), taught NextJS, FastAPI, and DevOps to BSc CS students.

EDUCATION

SASTRA DEEMED UNIVERSITY

Master of Computer Applications (MCA) - Chennai, India - 2021 - 2023